

Practical (?) Applications of Reflection

Jackie Kay

C++ Now

May 5th, 2017

jackieokay.com

"Narcissus", Caravaggio, circa 1597–1599





Alternative title image: "Metamorphosis of Narcissus", Salvador Dalí, 1937

What is reflection?

**Start not with the solution,
but the problem.**

Problem: code duplication due to hardcoding knowledge about the layout of a type and its members.

```
struct Box    { Point pos; float length; float width; };
struct Circle { Point pos; float radius; };

// Primitive shapes defined by rendering library
void draw(const Box&);
void draw(const Circle&);

template<typename T>
void draw(const T& shape) {
    // We know T is composed of only boxes and circles.
    // But how do we know how to access them?
}
```

One solution is costly at runtime:

```
struct Shape {  
    std::vector<Box> boxes;  
    std::vector<Circle> circles;  
};  
  
void draw(const Shape& shape) {  
    for (const auto& box : shape.boxes) {  
        draw(box);  
    }  
    for (const auto& circle : shape.circles) {  
        draw(circle);  
    }  
}
```

Another solution is costly for programmer time:

```
struct Capsule {  
    std::array<Circle, 2> ends;  
    Box middle;  
};  
  
void render(const Capsule& shape) {  
    for (const auto& circle : shape.ends) {  
        draw(circle);  
    }  
    draw(shape.middle);  
}  
  
// Imagine writing this for hundreds of objects  
// Now imagine maintaining it
```

How about an API that gets all the members for a generic type?

```
// Ideal syntax
```

```
template<typename T>
void render(const T& shape) {
    for... (auto& member : reflect(shape).members()) {
        if constexpr (member.renderable()) {
            render(member);
        }
    }
}
```


Let's call this "introspection".

C++ solutions for introspection

- Source code generation
 - Qt MOC
- "Adapt struct" macros
 - Fusion, Hana
- Compiler tooling
 - siplasplas
- Arcane secrets
 - POD Flat Reflection (formerly magic_get)

Qt MOC (Meta-Object compiler)

- Build step parses C++ code annotated with Qt macros
- Generates runtime metaobjects containing info like class name, methods, etc.
- Also generates signal/slot machinery, custom RTTI, properties

Adapt struct macros

Example lifted from [Boost Hana documentation](#):

```
struct Person {  
    BOOST_HANA_DEFINE_STRUCT(Person,  
        (std::string, name),  
        (std::string, last_name),  
        (int, age));  
};  
// or, if Person is externally defined  
BOOST_HANA_ADAPT_STRUCT(Person, name, last_name, age);
```

Adapt struct macros

```
Person presenter{"Jackie", "Kay", 24};
hana::for_each(presenter, [](auto pair) {
    std::cout << hana::to<char const*>(hana::first(pair))
               << ": " << hana::second(pair) << std::endl;
});
// name: Jackie
// last_name: Kay
// age: 24
```

How does it work?

- For each member, the macro:
- Generates generic set and get functions
- Stringizes the field name into constexpr string
- Uses the identifier name and type to get the member pointer
- Associates this into tuple of constexpr string, member pointer accessor pairs.
- Template specialization required for 1 member, 2 members, ... up to N members

POD Flat Reflection (pfr)

```
template<size_t I, typename T>
void print_members_recursive(const T& t) {
    if constexpr (I == pfr::tuple_size<T>{}) {
        return;
    } else {
        std::cout << pfr::flat_get<I>(t) << "\n";
        print_members_recursive<I + 1>(t);
    }
}

struct Person {
    const char* name;
    const char* last_name;
    int age;
};

Person presenter{"Jackie", "Kay", 24};
print_members_recursive<0>(presenter);
```

How does it work?

- C++17 version uses structured bindings and no macros
- That sounds much better... right?


```
// full namespaces and noexcept omitted for brevity
template <class T>
constexpr auto as_tuple_impl(T&& /*val*/, size_t_<0>) {
    return sequence_tuple::tuple<>{};
}

template <class T>
constexpr auto as_tuple_impl(T&& val, size_t_<1>) {
    auto& [a] = std::forward<T>(val);
    return detail::make_tuple_of_references(a);
}

template <class T>
constexpr auto as_tuple_impl(T&& val, size_t_<2>) {
    auto& [a,b] = std::forward<T>(val);
    return detail::make_tuple_of_references(a,b);
}

// etc...
```

```
template <class T>
constexpr auto as_tuple_impl(T&& val, size_t_<100>) {
    auto& [
        a,b,c,d,e,f,g,h,j,k,l,m,n,p,q,r,s,t,u,v,w,x,y,z,A,B,C,
        D,E,F,G,H,J,K,L,M,N,P,Q,R,S,T,U,V,W,X,Y,Z,aa,ab,ac,ad,
        ae,af,ag,ah,aj,ak,al,am,an,ap,aq,ar,as,at,au,av,aw,ax,
        ay,az,aA,aB,aC,aD,aE,aF,aG,aH,aJ,aK,aL,aM,aN,aP,aQ,aR,
        aS,aT,aU,aV,aW,aX,aY,aZ,ba,bb,bc,bd
    ] = std::forward<T>(val);

    return detail::make_tuple_of_references(
        a,b,c,d,e,f,g,h,j,k,l,m,n,p,q,r,s,t,u,v,w,x,y,z,A,B,C,
        D,E,F,G,H,J,K,L,M,N,P,Q,R,S,T,U,V,W,X,Y,Z,aa,ab,ac,ad,
        ae,af,ag,ah,aj,ak,al,am,an,ap,aq,ar,as,at,au,av,aw,ax,
        ay,az,aA,aB,aC,aD,aE,aF,aG,aH,aJ,aK,aL,aM,aN,aP,aQ,aR,
        aS,aT,aU,aV,aW,aX,aY,aZ,ba,bb,bc,bd
    );
}
```

siplasplas

Generate static introspection headers with libclang

```
class Foo { int i; };  
// generates in a separate header:  
template<>  
class Class<Foo>  
{  
public:  
    using Fields = typelist<Field<SourceInfo<  
        string<'i'>, // name  
        string<'f', 'o', 'o', '.', 'h'>, // file  
        4 // line  
    >,  
    decltype(&Foo::i), &Foo::i, // Pointer  
    >>;  
};
```

Other languages

- Python
 - Flexible and powerful, user-friendly syntax
 - Language design philosophy (every struct is a dict, weak typing) contributes to this
- Java
 - Uses runtime bytecode information, fairly extensive API
- C#
 - Arcane secrets

Proposed solutions for C++

reflexpr

- [P0194R3](#) by Matúš Chochlík and Axel Naumann
- I recommend P0578R0, ["Static Reflection in a Nutshell"](#) by Chochlík, Naumann and David Sankel
- [Clang implementation](#) and [documentation](#)

reflexpr: raw API

```
namespace meta = std::meta;  
using MetaInfo = reflexpr(Person);  
std::cout << "A " << meta::get_base_name_v<MetaInfo>  
           << " is made up of "  
           << meta::get_size<MetaInfo> << " things.\n";
```

operator\$/cpp3k

- [P0590R0](#) by Andrew Sutton and Herb Sutter
- [Clang implementation](#) also available
- Choice of \$ is controversial: supported as valid identifier as an extension in most compilers
 - But there's more to it than just the symbol

operator\$

```
namespace meta = cpp3k::meta::v1;  
constexpr auto info = $Person;  
std::cout << "A " << info.name() << " is made up of "  
           << info.member_variables().size()  
           << " things.\n";
```

**What can we do with
reflection?**

Warm-up: equality operators

```
template<typename T>
bool equal(const T& a, const T& b) {
    if constexpr (equality_comparable<T>()) {
        return a == b;
    } else if constexpr (iterable<T>()) {
        if (a.size() != b.size()) {
            return false;
        }
        for (int i = 0; i < a.size(); ++i) {
            if (!equal(a[i], b[i])) return false;
        }
        return true;
    } else {
        /* Time for reflection */
    }
}
```

With reflexpr

```
using MetaT = reflexpr(T);
static_assert(meta::Record<MetaT>,
    "Reached non-equality comparable leaf member.");
bool result = true;
meta::for_each<meta::get_data_members_m<MetaT>>(
    [&a, &b, &result](auto&& member) {
        using M = typename std::decay_t<decltype(member)>;
        constexpr auto p = meta::get_pointer<M>::value;
        result &= equal(a.*p, b.*p);
    }
);
return result;
```

With cpp3k

```
static_assert(refl::is_member_type<T>(),
    "Reached non-equality comparable leaf member.");
bool result = true;
meta::for_each($T.member_variables(),
    [&a, &b, &result](auto&& member){
        constexpr auto p = member.pointer();
        result &= equal(a.*p, b.*p);
    }
);
return result;
```

Observations

- Using types to represent metainfo values has more syntactic overhead, especially when using lambdas
- `for_each` requires valid initial state of result, doesn't allow early return/short circuiting
- What about private vs. public members?

Fold expressions

constexpr can express sequence as a parameter pack:

```
template<typename ...Pack>
struct compare_fold {
    static constexpr auto apply(const T& a, const T& b) {
        return (equal(a.*pointer<Pack>(), b.*pointer<Pack>())
            && ...);
    }
};
// ...
meta::unpack_sequence_t<meta::get_data_members_m<MetaT>,
    compare_fold>::apply(a, b);
```

Compiler optimizations on fold expression could provide short circuiting

Access specifiers

- `reflexpr` offers `get_data_members`, `get_public_data_members`, and `is_public/private` for individual member metainfo
- `cpp3k` provides `is_public/private` but no convenience function for all public members

Bypassing private is easier*

```
class A { int x = 42; };
struct B {
    int y;
    B(const A& a) {
        y = a.*meta::get_pointer<
            meta::get_element_m<
                meta::get_data_members_m<reflexpr(A)>, 0
            >>::value;
    }
};
// ...
A a;
assert(B(a).y == 42);
```

*see blog post by [Johannes Schaub](#)

My 2p on access specifiers

- Access specifiers are already "broken" in the language. They're still useful for interface design.
- Guideline: prefer to reflect on public data members for serialization, etc. and hide implementation details in private members, to avoid breaking compatibility between library versions
- Standardized API should include a convenience accessor for public member variables, to encourage this guideline

Serialization/deserialization

Louis Dionne's Meeting C++ 2016 keynote showed a JSON serializer using a value-semantics metaprogramming mini-library built on top of `reflexpr`.

How about deserialization?

```
template<typename T>  
auto deserialize(std::string_view& src, T& dst);
```

Requires matching a runtime string to a metainfo.

Linear matching

```
// input: string_view representing the string to parse  
// parsed_keys: all key strings for a JSON object  
// result: enum  
for (const auto& key : parsed_keys) {  
    const auto& value = parse_next_value(input);  
    meta::for_each($T.member_variables(),  
        [&dst, &key, &value, &result](auto&& m) {  
        if (key == m.name()) {  
            if (result = deserialize(value, dst.*m.pointer()));  
                result != deserialize_result::success) {  
                return;  
            }  
        }  
    }  
};  
}
```

Observations

- Same pattern of recursively applying an operation over all members of a struct as the equality operator example.
- String matching could be costly at runtime: if T has n members, $O(n)$ string comparisons per JSON key.

Program options

Program options: interface

Struct member name corresponds to command line flag and abbreviation.

```
struct ProgramOptions {  
    std::string filename;    // --filename, -f  
    int iterations = 100;    // --iterations, -i  
    bool help = false;       // --help, -h  
    float foo;               // --foo, -o  
};  
  
// Returns nullptr if there was a parsing error  
template<typename T>  
optional<T> parse(int argc, char** argv const);
```


Program options: outline

- Accumulate a Hana compile-time map with pairs: ("--member_name", metainfo)
- Take 'c', the first character in "member_name". If "-c" not in the map, add ("-c", "metainfo"), else check next character.
- `parse` iterates over `argv` and converts a runtime `const char*` to `constexpr` map key to retrieve metainfo
- Metainfo provides member pointer and type (needed to set member)

Program options: reflexpr

```
template<typename T> struct OptionsMap {
    using MetaT = reflexpr(T);
    template<typename... MetaFields>
    struct make_prefix_map {
        static constexpr auto helper() {
            return hana::fold(
                hana::make_tuple(MetaFields{}...),
                hana::make_map(), collect_flags
            );
        }
    };
    static constexpr auto prefix_map =
        meta::unpack_sequence_t<
            meta::get_data_members_m<MetaT>,
            make_prefix_map>::helper();
};
```

Program options: cpp3k

```
template<typename T> struct OptionsMap {
    static constexpr auto prefix_map = hana::fold(
        refl::adapt_to_hana($T.member_variables()),
        hana::make_map(),
        collect_flags
    );
};

/* "adapt_to_hana" internals */
template<typename T, size_t ...I>
constexpr auto adapt_to_hana_helper(const T& t,
    std::index_sequence<I...>&&) {
    return hana::make_tuple(
        cpp3k::meta::v1::cget<I>(t)...
    );
}
```

constexpr string parameters

- `get_base_name_v<reflexpr(T)>` returns a `const char[N]`
- `$T.name()` returns a `const char*`
- Can't pass these constexpr values to functions and use function parameters in a template context
- Furthermore, `hana::map` string keys must be `hana::string`'s

Stop-gap solution


Pass type of metainfo into a `get_name` function.

```
template<typename T, size_t... I>
constexpr auto helper(std::index_sequence<I...>) {
    return hana::string_c<T::name()[I]...>;
}
template<typename T>
constexpr auto get_name() {
    constexpr auto L = static_strlen(T::name());
    return helper<T>(std::make_index_sequence<L>{});
}
// Usage:
using T = std::decay_t<decltype(field)>;
hana::insert(prefix_map,
    hana::make_pair("--"_s + get_name<T>(), field));
```

Parsing and setting members

```
auto set(T& opt, const char* prefix, const char* val) {
    hana::for_each(hana::keys(prefix_map),
        [&options, &prefix, &v](const auto& key) {
            if (runtime_string_compare(key, prefix)) {
                constexpr auto info = hana::at_key(prefix_map,
                                                    decltype(key));
                constexpr auto p = info.pointer();
                using M = unreflect_member_t<T, decltype(info)>;
                opt.*p = boost::lexical_cast<M>(val, strlen(val));
            }
        }
    );
}
```

vs. boost::program_options



Observations

- Cconstexpr strings can be annoying with the current state of the language. We need a standard representation and utilities such as constexpr `strlen` and `strcmp`.
- In order to add metadata such as help strings or custom flags, my ideal syntax is user-defined attributes and reflection on attributes
- Same linear runtime string matching pattern appears. Can we do better?

constexpr string hashing

Idea: exploit our knowledge of the set of all member identifiers at compile time.

Implement a runtime perfect hash from N unique strings to N unique integers (no collisions).

Then match the runtime integer to a compile-time index and a callback.

In our reflection examples, the hash input is a member name, and the hash output is the compile-time index into the struct.

string hash implementation

A highly simplified version of the gperf algorithm

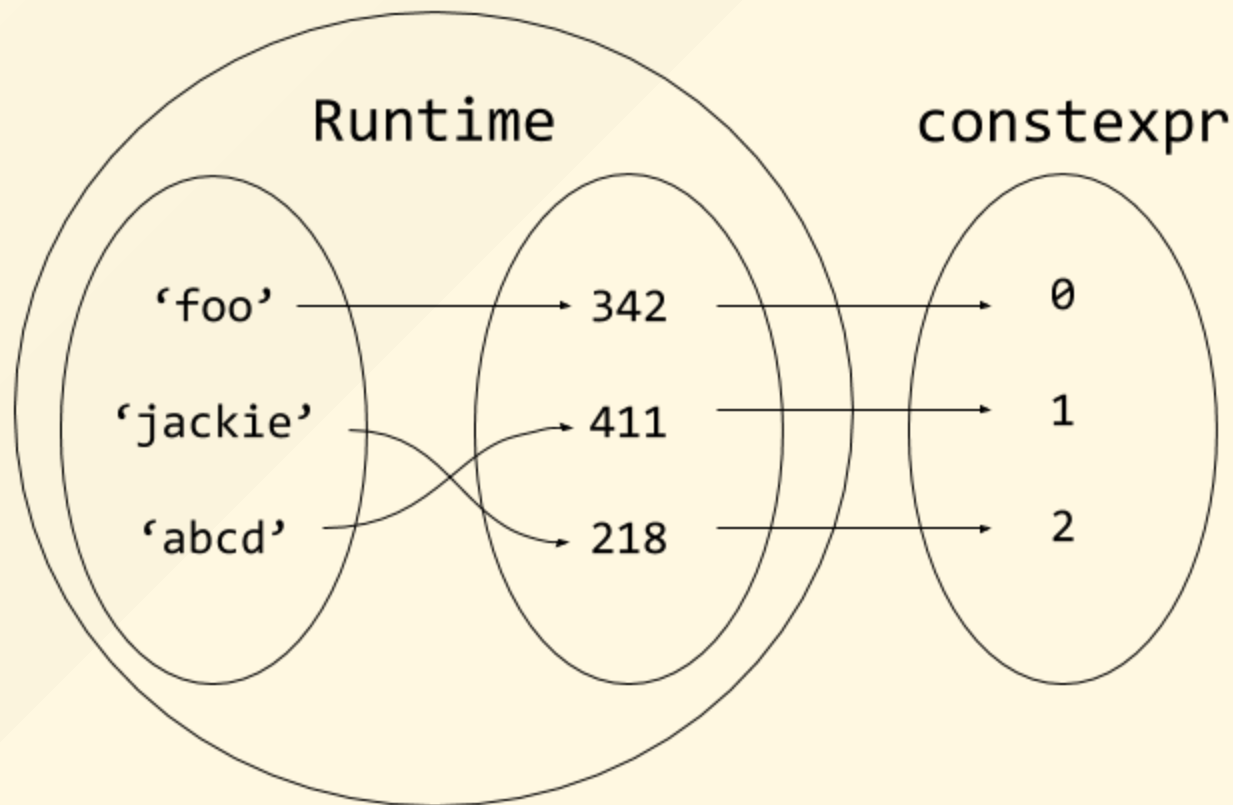
```
template<typename... Strings>
struct simple_string_hash {
    static constexpr auto positions =
        std::make_index_sequence<
            max_string_length(Strings{ }...)
        >{};
    static constexpr auto increments =
        std::make_index_sequence<
            decltype(positions)::size()
        >{};
    /* Implement runtime hash operators... */
};
```

string hash implementation

```
struct simple_string_hash {  
    template<size_t ...I>  
    static auto helper(  
        const char* key, const std::index_sequence<I...>&  
    )  
    {  
        return ((key[I] + sequence_at<I>(increments)) + ...);  
    }  
  
    auto operator()(const char* key) const {  
        return helper(key, positions) + strlen(key);  
    }  
};
```

Matching integer sets

String hash is perfect, but not sequential. Need another mapping



Matching integer sets

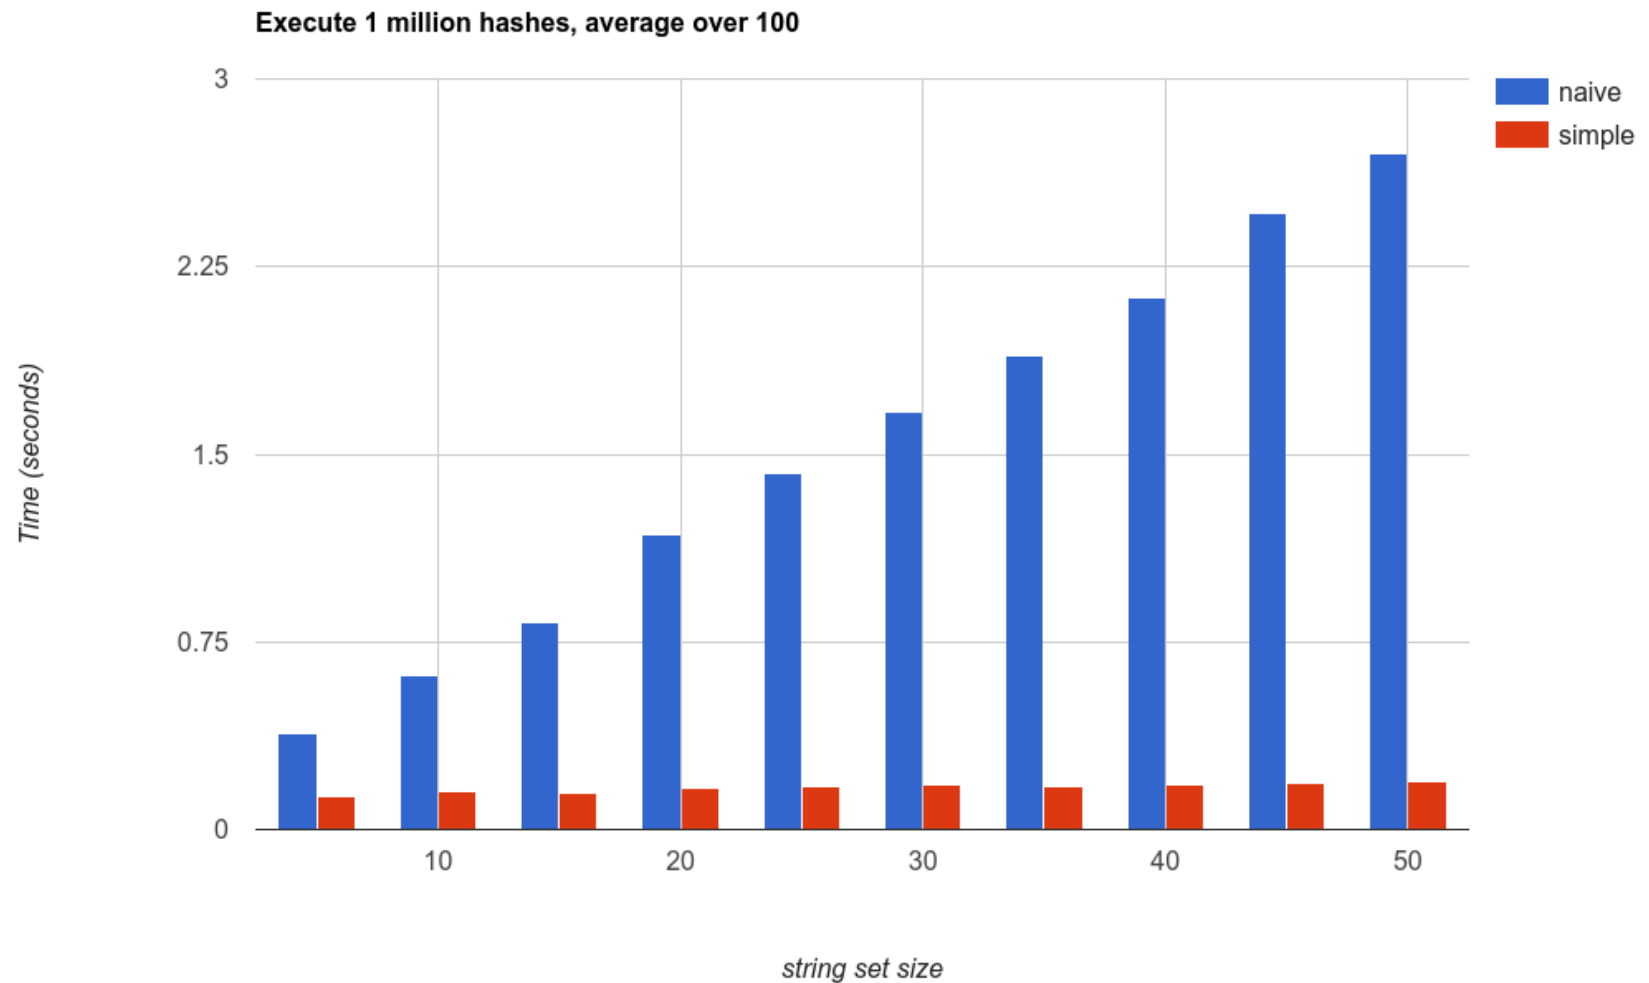
```
template<template<size_t> typename F, size_t ...Sequence>
struct recursive_switch_table {
    using IndexSeq = std::index_sequence<Sequence...>;
    template<size_t I, size_t Iterations>
    static auto apply(unsigned i) {
        switch(i) {
            case I:
                return F<map_to_index<I, Sequence...>()>{}();
            default:
                constexpr unsigned n = Iterations + 1;
                return apply<sequence_at<n>(IndexSeq{}), n>(i);
        }
    }
    void operator()(unsigned i) const {
        return apply<sequence_at<0>(IndexSeq{}), 0>(i);
    }
};
```

Benchmark

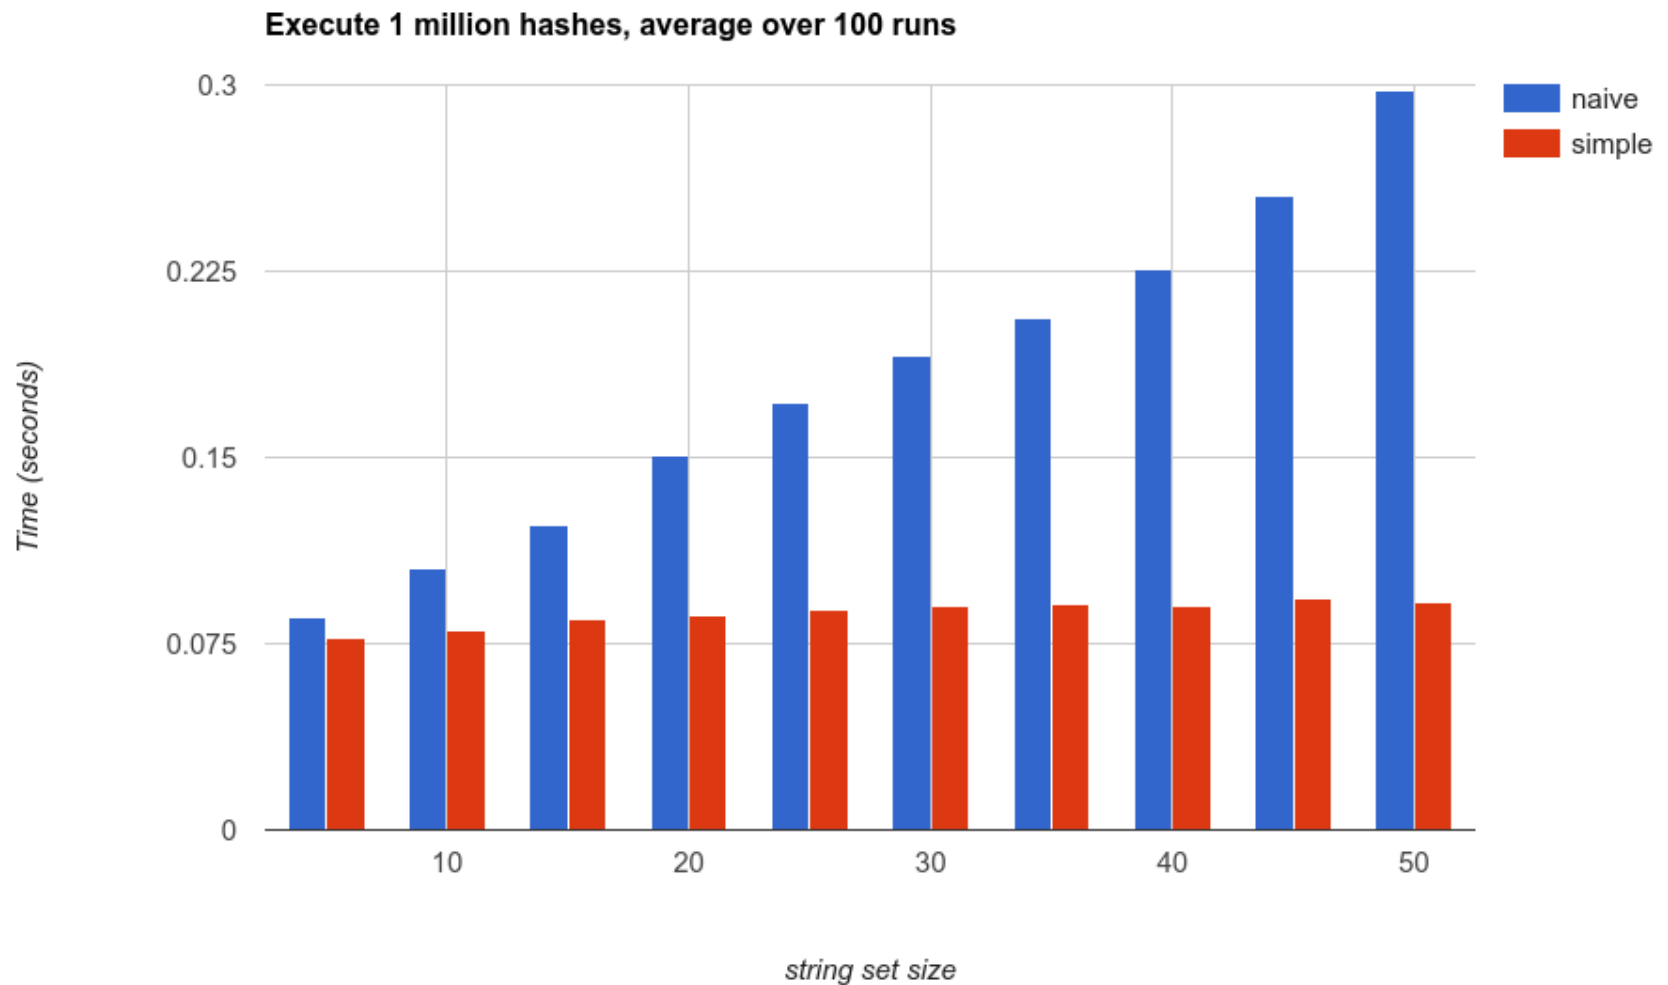
Compare to naive (linear) implementation:

```
// F: callback  
// StringSet: tuple of string literals  
template<template<size_t> typename F, typename StringSet,  
        size_t ...I>  
constexpr auto naive_string_hash(  
    StringSet&& set, std::index_sequence<I...>&&) {  
    return [set = std::move(set)](const char* input) {  
        ([&set, &input]() {  
            if (sl::equal(std::get<I>(string_set), input)) {  
                F<I>{}();  
            }  
        })(), ...);  
    };  
}
```

Results (-OO)



Results (-O3)



Caveats/observations

- Number of members in a struct rarely exceeds 10
- Interface to this hash could be improved
- Hash performance could be improved by "compressing" the representation (more work at compile time)
- But, with or without reflection, metaprogramming could become more "practical" with efficient runtime to compile-time mappings

Miscellaneous thoughts on performance

Compile times

- Type-based metaprogramming generally faster than value-based
 - Refer to metabেন.ch, Odin Holmes's blog and presentations
- This is the best justification for `reflexpr` style API
- $\text{priority}(\text{human time}) > \text{priority}(\text{compiler time})$?
Who decides?
- Is compiler speed a library issue, language issue, or QoI issue?

Shifting gears

Function reflection

Ever written cryptic boilerplate metafunctions to get the arity of a function or an invocable?

```
void f(int x, const std::string& name) {  
    // ...  
}  
  
// ...  
  
constexpr unsigned arity_f = $f.parameters().size();
```

(Note: \$ implementation requires name of function, doesn't support lambdas)

Reflection on overload sets

Warning: fictional syntax

```
void f(int x) {  
    // ...function body...  
}  
void f(float x) {  
    // ...function body...  
}  
  
// ...  
for... (const auto& overload : $f.overloads) {  
    // "int" in first iteration, "float" in the next  
    using ArgT = decltype(overload.parameters()[0])::type;  
}
```

What's missing?

**Introspection is only one
piece of the puzzle.**

**We need the ability to
manipulate identifiers at
compile-time.**

Manipulating identifiers

I'll call this operator `iddecl`, for "identifier declaration".

The idea is sometimes called `idreflexpr`, but I consider this mechanism orthogonal to reflection.

I don't have access to a compiler with it implemented, so let's dream a bit...

"iddecl" example

```
struct Person { int age; const char* name; };

template<typename T>
struct A {
    // Fictional syntax:
    // Parameter pack expansion for member declaration
    (get_member_types_m<reflexpr(T)>
     iddecl(get_data_members_m<reflexpr(T)>))...;
};

A<Person> a{24, "jackie"};
std::cout << a.name << " is " << a.age << " years old.\n";
```

"How is this different from inheritance?"

It's better

(where better == more powerful)

```
template<typename T>
struct POD {
    (get_member_types_m<reflexpr(T)>
        iddecl(filter<
            get_data_members_m<reflexpr(T)>, is_pod
        >)
    )...;
};
```

Mocking with iddecl

```
template<typename T>
struct Mock {
    using MetaT = reflexpr(T);
    (get_member_function_types_m<MetaT>
     iddecl(get_member_functions_m<MetaT>)
     (Args... args)
     {
         // Mock library interface collects
         return impl<get_member_functions_m<MetaT>>(args...);
     }
    )...;
};
```

With \$ and constexpr for (on steroids)

```
template<typename T>
struct Mock {
    constexpr {
        for...(const auto& function : $T.member_functions()) {
            using Function = decltype(function);
            Function::return_type
            iddecl(function)(Function::arguments&& args) {
                constexpr auto impl = impl_map.at(function);
                return std::apply(impl, args);
            }
        }
    }
};
```

Traits/interfaces ala [dyno](#)

```
struct DrawableInterface {  
    virtual void draw(std::ostream& out) = 0;  
};  
  
using drawable = concept_map<DrawableInterface>;  
  
struct Square {  
    void draw(std::ostream& out) const { out << "Square"; }  
};  
  
void f(drawable const& d) {  
    d.draw(std::cout);  
}
```


Traits/interfaces ala [dyno](#)

```
template<typename >
struct concept_map {
    constexpr {
        for...(const auto& function : $I.member_functions()) {
            using Function = decltype(function);
            Function::return_type
            iddecl(function)(Function::arguments&& args) {
                return poly_.virtual_(
                    stringize(function.name()))(poly, args);
            }
        }
    }
private:
    dyno::poly<Drawable> poly_;
};
```

And other customization points that suck less

Hint: watch Michał Dominiak's talk


Why is it missing right now?

There's no precedent in the language: can only manipulate identifiers with the preprocessor.

Can we use constexpr strings in an iddecl? How do you refer to the result of an iddecl? If you "clone" another struct's member with iddecl, do you maintain its default value, cv qualifiers, attributes, etc?

We need a strong vision going forward.

For that, I'll let this guy take the stage:



but you'll have to wait a few more months!

Acknowledgments

This presentation wouldn't exist without:

- Matúš Chochlík, Axel Naumann, and David Sankel for including me in discussions on reflection
- Extra thank you to Matúš for providing `reflexpr`
- Andrew Sutton for providing `operator$`
- Louis Dionne for Hana and string hashing ideas
- Vittorio Romeo for inspiring the program options example, feedback, and moral support

Presentation links

- github.com/jacquelinekay/reflection_experiments
- github.com/jacquelinekay/c++now2017
- github.com/jacquelinekay/dispatch